Source and claim audit — LLM Comparison Lab

Deck date: 2026-10-21. **Scope:** 17 slides on the current React/FastAPI app. No provider generation was sent for the deck. The older Flask/CLI prototype was excluded.

Claim	Checked basis	Boundary
Research frame	Chang et al. (2024); sources/ai_project.md	Original WHAT / WHERE / HOW diagram; bounded local adaptation.
Paper mapping	docs/RESEARCH_FOR_ARENA.md; docs/METHODOLOGY.md	HELM, IFEval, Dror and Zheng motivate or inspire local choices. No replication.
IFEval authorship	arXiv:2311.07911v1	Jeffrey Zhou is first author; Le Hou is last. Corrected the prompt's lead attribution.
Questions	services/datasets.py; README.md; docs/METHODOLOGY.md	30 authored sample items are not representative. Five-family draft is not held out.
Execution	services/experiments.py; worker/runner.py	Frozen design and attempts support audit; uncertain remote outcomes need acknowledgement.
Provider safeguards	docs/PROVIDER_CONFIGURATION.md; docs/VERIFICATION.md	Zen evidence rechecked Oct 21; local expiry Oct 22 00:00 UTC. Catalog selection does not prove authenticated generation.
Scoring and interval	evaluation/scoring.py; services/analysis.py	Local checks and 10,000-draw family bootstrap. Convenience sampling limits inference.
Results evidence	docs/screenshots/results-desktop.png; AGENTS.md	Screenshot is synthetic. Separate 20/20 live smoke is a 10-item, two-OpenRouter-route convenience run.
Human / AI judge	frontend/src/HumanReview.tsx; docs/AI_JUDGE_DESIGN.md	Local review exists; independent raters unverified. AI judge remains design only.

Bibliography: five core papers (Chang, Liang, Zhou, Dror, Zheng) and three contextual 2025 surveys (Cao, Mohammadi, Ni). Links appear in slides 16–17. The contextual surveys are not implemented methods.

Known documentation conflict: AGENTS.md and part of README.md retain outdated Zen-key wording; current provider configuration and implementation supersede it.